

#PySpark Task Set - Part 3

#HR & Workforce Analytics, DataFrame APIs, Joins, SQL, Date Logic, Aggregation, UDFs, Views

```
from pyspark.sql import SparkSession
import pandas as pd
from io import StringIO
```

```
spark=SparkSession.builder.appName("HR Workforce Analytics").getOrCreate()
```

```
print(spark)
```

```
↳ <pyspark.sql.session.SparkSession object at 0x7d5a6e281c10>
```

#Data

#Sample CSV and JSON files

```
employees_csv = """EmpID,Name,Department,JoinDate,Salary,ManagerID
```

```
1,Anita,HR,2021-05-01,55000,
2,Raj,Engineering,2020-03-15,80000,1
3,Simran,Engineering,2022-07-10,75000,1
4,Aamir,Marketing,2019-11-20,60000,1
5,Nisha,HR,2023-01-05,50000,1
"""
```

```
attendance_csv = """EmpID,Date,Status
```

```
1,2024-04-01,Present
1,2024-04-02,Present
2,2024-04-01,Absent
2,2024-04-02,Present
3,2024-04-01,Present
3,2024-04-02,Present
4,2024-04-01,Absent
4,2024-04-02,Absent
5,2024-04-01,Present
5,2024-04-02,Present
"""
```

```
bonuses_json = """
```

```
[
{"EmpID": 1, "Year": 2023, "Bonus": 5000},
{"EmpID": 2, "Year": 2023, "Bonus": 7000},
{"EmpID": 3, "Year": 2023, "Bonus": 6500},
{"EmpID": 4, "Year": 2023, "Bonus": 6000},
{"EmpID": 5, "Year": 2023, "Bonus": 4000}
]
"""
```

Saving

```
with open("employees.csv", "w") as f:
    f.write(employees_csv)
```

```
with open("attendance.csv", "w") as f:
    f.write(attendance_csv)
```

```
with open("bonuses.json", "w") as f:
    f.write(bonuses_json)
```

#Task 1- Ingestion & Exploration

```
from pyspark.sql.functions import *
```

```
employees = spark.read.csv("employees.csv", header=True, inferSchema=True)
attendance = spark.read.csv("attendance.csv", header=True, inferSchema=True)
bonuses = spark.read.json("bonuses.json")
```

```
employees.printSchema()
attendance.printSchema()
bonuses.printSchema()
```

```
employees.show()
attendance.show()
bonuses.show()
```

Distinct departments

```
distinct_dept = employees.select("Department").distinct()
```

```
distinct_dept.count()
distinct_dept.show()

|-- ManagerID: integer (nullable = true)
```

```
root
|-- EmpID: integer (nullable = true)
|-- Date: date (nullable = true)
|-- Status: string (nullable = true)
```

```
root
|-- Bonus: long (nullable = true)
|-- EmpID: long (nullable = true)
|-- Year: long (nullable = true)
|-- _corrupt_record: string (nullable = true)
```

```
+-----+-----+-----+-----+
|EmpID|  Name| Department|  JoinDate|Salary|ManagerID|
+-----+-----+-----+-----+
|  1| Anita|      HR|2021-05-01| 55000|      NULL|
|  2|   Raj|Engineering|2020-03-15| 80000|         1|
|  3|Simran|Engineering|2022-07-10| 75000|         1|
|  4|  Aamir|Marketing|2019-11-20| 60000|         1|
|  5|  Nisha|      HR|2023-01-05| 50000|         1|
+-----+-----+-----+-----+
```

```
+-----+-----+-----+
|EmpID|      Date| Status|
+-----+-----+-----+
|  1|2024-04-01|Present|
|  1|2024-04-02|Present|
|  2|2024-04-01| Absent|
|  2|2024-04-02|Present|
|  3|2024-04-01|Present|
|  3|2024-04-02|Present|
|  4|2024-04-01| Absent|
|  4|2024-04-02| Absent|
|  5|2024-04-01|Present|
|  5|2024-04-02|Present|
+-----+-----+-----+
```

```
+-----+-----+-----+
|Bonus|EmpID|Year|_corrupt_record|
+-----+-----+-----+
| NULL| NULL|NULL|                [|
| 5000|   1|2023|                NULL|
| 7000|   2|2023|                NULL|
| 6500|   3|2023|                NULL|
| 6000|   4|2023|                NULL|
| 4000|   5|2023|                NULL|
| NULL| NULL|NULL|                ]|
+-----+-----+-----+
```

```
+-----+
| Department|
+-----+
|Engineering|
|      HR|
|Marketing|
+-----+
```

#Task 2-Data Frame Operations

```
from pyspark.sql.functions import datediff, current_date, round
```

```
#Add a column TenureYears using datediff() and round()
emp_df = employees.withColumn("TenureYears", round(datediff(current_date(), col("JoinDate")) / 365, 2))
emp_df.show()
```

```
#Calculate TotalCompensation = Salary + Bonus .
emp_bonus_df = emp_df.join(bonuses, "EmpID")
emp_bonus_df = emp_bonus_df.withColumn("TotalCompensation", col("Salary") + col("Bonus"))
emp_bonus_df.show()
```

```
#Filter employees with more than 2 years in the company.
filtered_emp = emp_bonus_df.filter(col("TenureYears") > 2)
filtered_emp.show()
```

```
#Show employees who report to a manager ( ManagerID is not null ).
emp_bonus_df.filter(col("ManagerID").isNotNull()).show()
```



EmpID	Name	Department	JoinDate	Salary	ManagerID	TenureYears
1	Anita	HR	2021-05-01	55000	NULL	4.11
2	Raj	Engineering	2020-03-15	80000	1	5.24
3	Simran	Engineering	2022-07-10	75000	1	2.92
4	Aamir	Marketing	2019-11-20	60000	1	5.56
5	Nisha	HR	2023-01-05	50000	1	2.43

EmpID	Name	Department	JoinDate	Salary	ManagerID	TenureYears	Bonus	Year	_corrupt_record	TotalCompensation
1	Anita	HR	2021-05-01	55000	NULL	4.11	5000	2023	NULL	60000
2	Raj	Engineering	2020-03-15	80000	1	5.24	7000	2023	NULL	87000
3	Simran	Engineering	2022-07-10	75000	1	2.92	6500	2023	NULL	81500
4	Aamir	Marketing	2019-11-20	60000	1	5.56	6000	2023	NULL	66000
5	Nisha	HR	2023-01-05	50000	1	2.43	4000	2023	NULL	54000

EmpID	Name	Department	JoinDate	Salary	ManagerID	TenureYears	Bonus	Year	_corrupt_record	TotalCompensation
1	Anita	HR	2021-05-01	55000	NULL	4.11	5000	2023	NULL	60000
2	Raj	Engineering	2020-03-15	80000	1	5.24	7000	2023	NULL	87000
3	Simran	Engineering	2022-07-10	75000	1	2.92	6500	2023	NULL	81500
4	Aamir	Marketing	2019-11-20	60000	1	5.56	6000	2023	NULL	66000
5	Nisha	HR	2023-01-05	50000	1	2.43	4000	2023	NULL	54000

EmpID	Name	Department	JoinDate	Salary	ManagerID	TenureYears	Bonus	Year	_corrupt_record	TotalCompensation
2	Raj	Engineering	2020-03-15	80000	1	5.24	7000	2023	NULL	87000
3	Simran	Engineering	2022-07-10	75000	1	2.92	6500	2023	NULL	81500
4	Aamir	Marketing	2019-11-20	60000	1	5.56	6000	2023	NULL	66000
5	Nisha	HR	2023-01-05	50000	1	2.43	4000	2023	NULL	54000

#Task 3- Aggregation

#Avg salary per department

```
avg_salary_dept = employees.groupBy("Department").agg(avg("Salary"))
avg_salary_dept.show()
```

#No. of Employees under each manager

```
no_of_emp_manager = employees.groupBy("ManagerID").agg(count("*").alias("TeamSize"))
no_of_emp_manager.show()
```

#Count of absences per employee

```
count_of_absence = attendance.filter(col("Status") == "Absent") \
    .groupBy("EmpID").count().withColumnRenamed("count", "AbsenceCount")
count_of_absence.show()
```



Department	avg(Salary)
Engineering	77500.0
HR	52500.0
Marketing	60000.0

ManagerID	TeamSize
NULL	1
1	4

EmpID	AbsenceCount
4	2
2	1

#4. Joins

```
total_days = attendance.groupBy("EmpID").count().withColumnRenamed("count", "TotalDays")
present_days = attendance.filter(col("Status") == "Present") \
```

```

.groupBy("EmpID").count().withColumnRenamed("count", "PresentDays")

attendance_rate = total_days.join(present_days, "EmpID") \
    .withColumn("AttendancePercent", round(col("PresentDays") / col("TotalDays") * 100, 2))

#Join employees and attendance → Get attendance % (Present days / Total days).
emp_attendance = employees.join(attendance_rate, "EmpID", "left")
emp_attendance.select("EmpID", "Name", "AttendancePercent").show()

#Join employees and bonuses → Show top 3 employees by TotalCompensation.
emp_bonus_totalcomp = emp_bonus_df.orderBy(col("TotalCompensation").desc()).select("EmpID", "Name", "TotalCompensation")
emp_bonus_totalcomp.show(3)

#Multi-level join: employees + bonuses + attendance .
full_df = employees.join(bonuses, "EmpID").join(attendance_rate, "EmpID", "left")
full_df.show()

```



```

+-----+-----+-----+
|EmpID|  Name|AttendancePercent|
+-----+-----+-----+
| 1| Anita|          100.0|
| 2|  Raj|           50.0|
| 3|Simran|          100.0|
| 4| Aamir|           NULL|
| 5| Nisha|          100.0|
+-----+-----+-----+

```

```

+-----+-----+-----+
|EmpID|  Name|TotalCompensation|
+-----+-----+-----+
| 2|  Raj|          87000|
| 3|Simran|          81500|
| 4| Aamir|          66000|
+-----+-----+-----+

```

only showing top 3 rows

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|EmpID|  Name| Department|  JoinDate|Salary|ManagerID|Bonus|Year|_corrupt_record|TotalDays|PresentDays|AttendancePercent|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1| Anita|      HR|2021-05-01| 55000|    NULL| 5000|2023|          NULL|      2|      2|          100.0|
| 2|  Raj|Engineering|2020-03-15| 80000|      1| 7000|2023|          NULL|      2|      1|           50.0|
| 3|Simran|Engineering|2022-07-10| 75000|      1| 6500|2023|          NULL|      2|      2|          100.0|
| 4| Aamir|Marketing|2019-11-20| 60000|      1| 6000|2023|          NULL|    NULL|    NULL|           NULL|
| 5| Nisha|      HR|2023-01-05| 50000|      1| 4000|2023|          NULL|      2|      2|          100.0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

#5. String & Date Functions

#Extract year and month from JoinDate .

```

employees = employees.withColumn("JoinYear", year("JoinDate")) \
    .withColumn("JoinMonth", month("JoinDate"))
employees.show()

```

#Mask employee names using regex.

```

employees = employees.withColumn("MaskedName", regexp_replace("Name", "[aeiouAEIOU]", "*"))
employees.show()

```

#Use substring() to create EmpCode like "EMP001".

```

employees = employees.withColumn("EmpCode", format_string("EMP%03d", col("EmpID")))
employees.select("EmpID", "Name", "EmpCode", "MaskedName").show()

```



```

+-----+-----+-----+-----+-----+-----+-----+-----+
|EmpID|  Name| Department|  JoinDate|Salary|ManagerID|JoinYear|JoinMonth|
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1| Anita|      HR|2021-05-01| 55000|    NULL|  2021|      5|
| 2|  Raj|Engineering|2020-03-15| 80000|      1|  2020|      3|
| 3|Simran|Engineering|2022-07-10| 75000|      1|  2022|      7|
| 4| Aamir|Marketing|2019-11-20| 60000|      1|  2019|     11|
| 5| Nisha|      HR|2023-01-05| 50000|      1|  2023|      1|
+-----+-----+-----+-----+-----+-----+-----+-----+

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|EmpID|  Name| Department|  JoinDate|Salary|ManagerID|JoinYear|JoinMonth|MaskedName|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1| Anita|      HR|2021-05-01| 55000|    NULL|  2021|      5|  *n*t*|
| 2|  Raj|Engineering|2020-03-15| 80000|      1|  2020|      3|  R*j|
| 3|Simran|Engineering|2022-07-10| 75000|      1|  2022|      7|  S*mr*n|
| 4| Aamir|Marketing|2019-11-20| 60000|      1|  2019|     11|  **m*r|
| 5| Nisha|      HR|2023-01-05| 50000|      1|  2023|      1|  N*sh*|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

```

+-----+-----+-----+-----+
| EmpID | Name | EmpCode | MaskedName |
+-----+-----+-----+-----+
| 1 | Anita | EMP001 | *n*t* |
| 2 | Raj | EMP002 | R*j |
| 3 | Simran | EMP003 | S*mr*n |
| 4 | Aamir | EMP004 | **m*r |
| 5 | Nisha | EMP005 | N*sh* |
+-----+-----+-----+-----+

```

#6. Conditional & Null Handling

#Use when/otherwise to label performance:

"High" if Bonus > 6000

"Medium" if 4000-6000

"Low" otherwise

```

emp_bonus_df = emp_bonus_df.withColumn(
    "Performance",
    when(col("Bonus") > 6000, "High")
    .when((col("Bonus") > 4000) & (col("Bonus") <= 6000), "Medium")
    .otherwise("Low")
)
emp_bonus_df.show()

```

#Handle missing ManagerID using fillna("No Manager") .

```

employees_filled = employees.fillna({"ManagerID": "No Manager"})
employees_filled.show()

```

```

➡ +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| EmpID | Name | Department | JoinDate | Salary | ManagerID | TenureYears | Bonus | Year | _corrupt_record | TotalCompensation | Performance |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | Anita | HR | 2021-05-01 | 55000 | NULL | 4.11 | 5000 | 2023 | NULL | 60000 | Medium |
| 2 | Raj | Engineering | 2020-03-15 | 80000 | 1 | 5.24 | 7000 | 2023 | NULL | 87000 | High |
| 3 | Simran | Engineering | 2022-07-10 | 75000 | 1 | 2.92 | 6500 | 2023 | NULL | 81500 | High |
| 4 | Aamir | Marketing | 2019-11-20 | 60000 | 1 | 5.56 | 6000 | 2023 | NULL | 66000 | Medium |
| 5 | Nisha | HR | 2023-01-05 | 50000 | 1 | 2.43 | 4000 | 2023 | NULL | 54000 | Low |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| EmpID | Name | Department | JoinDate | Salary | ManagerID | JoinYear | JoinMonth | MaskedName | EmpCode |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | Anita | HR | 2021-05-01 | 55000 | NULL | 2021 | 5 | *n*t* | EMP001 |
| 2 | Raj | Engineering | 2020-03-15 | 80000 | 1 | 2020 | 3 | R*j | EMP002 |
| 3 | Simran | Engineering | 2022-07-10 | 75000 | 1 | 2022 | 7 | S*mr*n | EMP003 |
| 4 | Aamir | Marketing | 2019-11-20 | 60000 | 1 | 2019 | 11 | **m*r | EMP004 |
| 5 | Nisha | HR | 2023-01-05 | 50000 | 1 | 2023 | 1 | N*sh* | EMP005 |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

#7. Spark SQL

Create Database hr

```
spark.sql("CREATE DATABASE IF NOT EXISTS hr")
```

```
spark.catalog.setCurrentDatabase("hr")
```

Save dataframes as tables

```
employees.write.mode("overwrite").saveAsTable("employees")
```

```
attendance.write.mode("overwrite").saveAsTable("attendance")
```

```
bonuses.write.mode("overwrite").saveAsTable("bonuses")
```

SQL Queries

#Top paid employee in each department.

```

spark.sql("""
    SELECT Department, Name, Salary
    FROM employees
    WHERE (Department, Salary) IN (
        SELECT Department, MAX(Salary)
        FROM employees
        GROUP BY Department
    )
""").show()

```

#Attendance rate by department.

```

spark.sql("""
    SELECT
        e.Department,
        ROUND(AVG(CASE WHEN a.Status = 'Present' THEN 1 ELSE 0 END) * 100, 2) AS AttendanceRate
    FROM employees e
    JOIN attendance a ON e.EmpID = a.EmpID
""")

```

```
GROUP BY e.Department
""").show()

#Employees joined after 2021 with salary >
70,000.
spark.sql("""
SELECT * FROM employees
WHERE JoinDate > '2021-01-01' AND Salary > 70000
""").show()
```

```
↗
+-----+-----+
| Department| Name|Salary|
+-----+-----+
|           | HR|Anita| 55000|
|Engineering| Raj| 80000|
|Marketing|Aamir| 60000|
+-----+-----+
```

```
+-----+-----+
| Department|AttendanceRate|
+-----+-----+
|Engineering|           75.0|
|           | HR           100.0|
|Marketing|           0.0|
+-----+-----+
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|EmpID| Name| Department| JoinDate|Salary|ManagerID|JoinYear|JoinMonth|MaskedName|EmpCode|
+-----+-----+-----+-----+-----+-----+-----+-----+
|      3|Simran|Engineering|2022-07-10| 75000|          1|    2022|          7| S*mr*n| EMP003|
+-----+-----+-----+-----+-----+-----+-----+-----+
```

```
#Task 8- Advanced(Optional)
from pyspark.sql.functions import udf
from pyspark.sql.types import StringType

# UDF for Department classification
def classify_dept(dept):
    return "Tech" if dept in ["Engineering"] else "Non-Tech"

classify_udf = udf(classify_dept, StringType())

classified_df = employees.withColumn("DeptType", classify_udf("Department"))
classified_df.select("EmpID", "Department", "DeptType").show()

# Assuming emp_attendance DataFrame is already created
emp_attendance.createOrReplaceTempView("emp_attendance_summary")

summary_df = spark.sql("SELECT * FROM emp_attendance_summary")
summary_df.show(truncate=False)

# Save as parquet partitioned by Department
summary_df.write.mode("overwrite").partitionBy("Department").parquet("emp_attendance_summary.parquet")

# Read back parquet to confirm and show
parquet_df = spark.read.parquet("emp_attendance_summary.parquet")
parquet_df.show(truncate=False)
```

```
↗
+-----+-----+-----+
|EmpID| Department|DeptType|
+-----+-----+-----+
|      1|           |HR|Non-Tech|
|      2|Engineering|Tech|
|      3|Engineering|Tech|
|      4|Marketing|Non-Tech|
|      5|           |HR|Non-Tech|
+-----+-----+-----+
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|EmpID|Name| Department|JoinDate|Salary|ManagerID|TotalDays|PresentDays|AttendancePercent|
+-----+-----+-----+-----+-----+-----+-----+-----+
|1|Anita|HR|2021-05-01|55000|NULL|2|2|100.0|
|2|Raj|Engineering|2020-03-15|80000|1|2|1|50.0|
|3|Simran|Engineering|2022-07-10|75000|1|2|2|100.0|
|4|Aamir|Marketing|2019-11-20|60000|1|NULL|NULL|NULL|
|5|Nisha|HR|2023-01-05|50000|1|2|2|100.0|
+-----+-----+-----+-----+-----+-----+-----+-----+
```

EmpID	Name	JoinDate	Salary	ManagerID	TotalDays	PresentDays	AttendancePercent	Department
1	Anita	2021-05-01	55000	NULL	2	2	100.0	HR
5	Nisha	2023-01-05	50000	1	2	2	100.0	HR
2	Raj	2020-03-15	80000	1	2	1	50.0	Engineering
3	Simran	2022-07-10	75000	1	2	2	100.0	Engineering
4	Aamir	2019-11-20	60000	1	NULL	NULL	NULL	Marketing